



Risk Radar

AI project-health monitor — flags risks, dependencies & blockers early

Challenge #1 x #3

Risk Radar x Status Summarizer
Dummy data · "Project Atlas"

 LIVE DEMO (INTERACTIVE)

arundhatimahapatro-holymuse.github.io/risk-radar/risk-radar.html


 PUBLIC REPOSITORY

github.com/arundhatimahapatro-holymuse/risk-radar

100-WORD SUMMARY

Risk Radar ingests messy, multi-source project updates — standups, Slack, Jira, status emails, PR notes — and turns them into an early-warning dashboard. A single engineered Claude prompt extracts and *separates* three things teams constantly conflate: **risks** (could go wrong), **dependencies** (X waits on Y), and **blockers** (stopped now), scoring each by impact × likelihood and citing the exact source phrase as evidence. It outputs strict JSON — overall RED/AMBER/GREEN health, a risk register, a cross-team dependency map, and prioritized next steps — that renders live in the dashboard. The browser prototype runs a heuristic mirror of the prompt, so it demos instantly with no API key.

THE PROTOTYPE



Live dashboard reading the "Project Atlas" updates → RED health 34/100, 2 blockers, 7 risks, 5 dependencies, a category risk-surface radar, and prioritized next steps.

HOW THE WORKFLOW RUNS

1 · Ingest	Pull last-24h updates from connectors (Slack, Jira, email, standup bot).
2 · Analyze	Feed raw text to the Risk Radar Claude prompt → strict JSON.
3 · Radar	Render dashboard, push a digest that @mentions owners of RED items, log score for week-over-week trend.

WHAT IT SEPARATES

Risk

Could go wrong.
Scored impact × likelihood.

Dependency

X waits on Y — a cross-team handoff.

Blocker

Work stopped right now.

THE ENGINEERED PROMPT (EXCERPT)

```
# SYSTEM – Risk Radar, a project-delivery risk analyst
- Be skeptical. Hedged language ("might slip", "still waiting") is an EARLY SIGNAL.
- Separate RISKS (could go wrong) · DEPENDENCIES (X needs Y) · BLOCKERS (stopped now).
- Score every risk: severity = impact × likelihood (High / Med / Low).
- Cite the exact source phrase as `evidence` for every item. Never invent facts.
- Health: RED if any active blocker or High risk on critical path; AMBER if medium; GREEN if on track.

# OUTPUT – strict JSON
{ "health": {status, score, rationale, trend},
  "risks": [{title, category, severity, likelihood, impact, evidence, source}],
  "dependencies": [{from, to, status, note}],
  "blockers": [{title, owner, since, evidence}],
  "next_steps": [{action, owner, priority}] }
```

Role + mission: "surface risk early"

Hedged language = signal

severity = impact × likelihood

evidence = exact source phrase (anti-hallucination)

deterministic RED/AMBER/GREEN rules

strict JSON → dashboard / Slack digest

Risk Radar · AI Workflow challenge submission · self-created dummy data, no confidential information.

Prototype: single-file HTML · Production: swap heuristic for Claude (Opus 4.8) prompt.